

Lecture 23 — Password Cracking, Bitcoin Mining, LLMs

Patrick Lam & Jeff Zarnett

`patrick.lam@uwaterloo.ca, jzarnett@uwaterloo.ca`

Department of Electrical and Computer Engineering
University of Waterloo

September 27, 2026

Part I

Password Cracking

scrypt is the algorithm behind Dogecoin.

The reference:

Colin Percival, “Stronger Key Derivation via Sequential Memory-Hard Functions”.

Presented at BSDCan’09, May 2009.

<http://www.tarsnap.com/scrypt.html>


- **not** plaintext!
- hashed and salted

One-way function:

- $x \mapsto f(x)$ easy to compute; but
- $f(x) \overset{?}{\mapsto} x$ hard to reverse.


Examples: SHA1, script.

Perhaps passwords have already been leaked!



The first thing is to try really common passwords.

You just might get a hit!



“All too easy.”

How can we reverse the hash function?

- Brute force.

GPUs (or custom hardware) are good at that!

The Arms Race: Making Cracking Difficult


Historically: force repeated iterations of hashing.

Main idea behind scrypt (hence Dogecoin):

- make hashing expensive in time and space.

Implication: require more circuitry to break passwords.
Increases both # of operations and cost of brute-forcing.

Of course, there's always this form of cracking:



(Source: xkcd 538)

Formalizing “expensive in time and space”

Definition

A memory-hard algorithm on a Random Access Machine is an algorithm which uses $S(n)$ space and $T(n)$ operations, where $S(n) \in \Omega(T(n)^{1-\epsilon})$.

Such algorithms are expensive to implement in either hardware or software.

Next, add a quantifier:
move from particular algorithms to underlying functions.

A sequential memory-hard function is one where:

- the fastest sequential algorithm is memory-hard; and
- it is impossible for a parallel algorithm to asymptotically achieve lower cost.

Exhibit. ReMix is a concrete example of a sequential memory hard function.

The script paper concludes with an example of a more realistic (cache-aware) model and a function in that context, BlockMix.


If designed well, hash functions aren't easy to brute force.

What if we remembered some previous work?


It isn't practical, or possible, to store the hashes for every possible plaintext.

The rainbow table is a compromise between speed and space.

The “reduction” function maps hashes to plaintext:



Example: $123456 \rightarrow d41d8cd98f00b204e9800998ecf8427e \rightarrow 418980$



We should do this to develop some number of chains.

This is the sort of task you could do with a GPU, because they can do a reduction relatively efficiently.

Or, y'know, just download them

Once we have those developed for a specific input set and hash function, they can be re-used forever.

You do not even need to make them yourself anymore (if you don't want) because you can download them on the internet... they are not hard to find.

They are large, yes, but in the 25–900 GB range.



I mean, Fallout 76 had a day one patch of 52 GB, and it was a disaster of a game.

- 1 Look for the hash in the list of final hashes; if there, we are done
- 2 If it's not there, reduce the hash into another plaintext and hash the new plaintext
- 3 Go back to step 1
- 4 If the hash matches a final hash, the chain with the match contains the original hash
- 5 Having identified the correct chain, we can start at the beginning of the chain with the starting plaintext and hash, check to see if we are successful (if so, we are done); if not, reduce and try the next plaintext.

Like generation, checking the tables can also be done efficiently by the GPU.

Some numbers from

<http://www.cryptohaze.com/gpurainbowcracker.php.html>:

- Table generation on a GTX295 core for MD5 proceeds at around 430M links/sec.
- Cracking a password 'K#n&r4Z':
real: 1m51.962s, user: 1m4.740s. sys: 0m15.320s

Yikes.

Part II

Bitcoin Mining

The World is Not Enough



Carbon Footprint

95.86 Mt CO₂



Comparable to the carbon footprint of Uzbekistan.

Electrical Energy

171.87 TWh



Comparable to the power consumption of Poland.

Electronic Waste

5.31 kt



Comparable to the small IT equipment waste of the Netherlands.

Fresh Water Consumption

2,709 GL



Comparable to the total water use of Switzerland.

<https://digiconomist.net/bitcoin-energy-consumption> As of July 2024

Basis: SHA-256

Started with CPU... then GPU... then what?



Custom hardware: optimize exactly what you need.

First FPGA, then ASIC...

We've uncovered why you should not mine Bitcoin.

Not cost efficient; way too much competition.


Not a hardware course, but this is the logical extension of GPU...

Part III

Large Language Models

In November of 2022, OpenAI introduced ChatGPT to the world...









I enjoy rephrasing
my question 15
different ways until
I get the answer I want



I'm a
Prompt Engineer

The "buy my course" starter pack



Is banned from PayPal



Doesn't want you to know that that his money really comes from selling his course



"Not saturated"



"Are you tired of working a 9-5?"

"I'm a mentor"



"college is a waste of money"



Usually in his late teens



"Five easy steps"



"you gotta find a Niche"

Sets up an "us vs them" dynamic

Such large language models have existed before, but ChatGPT ended up a hit because it's pretty good at being “conversational”.

This is within the ambit of Natural Language Processing (NLP).

Just because it gives you an answer, doesn't mean the answer is correct or true.



AI is subject to **hallucinations**.

Legally and professionally, the engineer is responsible for understanding how a given tool works & verifying the output is reasonable & correct.



Part of what makes the GPT-3 and GPT-4 models better at producing output that matches our expectations is that it relies on pre-training.

This course is not one on neural networks, large language models, AI, or similar.

But performance is relevant here!

One factor that matters in how good a model is: parameters.

Bigger is *usually* better...

But requires more computational and memory resources.

We can maybe tune some options!

This section is based on a guide from “Hugging Face”.



Hugging Face

You may have guessed by the placement of this topic in the course material that the GPU is the right choice for how to generate or train a large language model.

Just... uh... don't confuse Hugging Face with Facehugger...



In this case we're talking about Transformers.

There are three main groups of optimizations that it does:

- Tensor Contractions
- Statistical Normalizations
- Element-Wise Operators

We also need to consider what's in memory.

We can focus on how to generate a model that gives answers quickly...

Or we can focus on how to generate or train the model quickly – this will be our focus.

Use more space to reduce CPU usage, optimize for common cases, speculate...

Some of these are more fun than others: given a particular question, can you guess what the followup might be?



Why would we customize some LLM?

Don't send your data to OpenAI...

Specialize for your workload.

There are a lot of knobs we can tweak with respect to training the model.

We'll explore the configuration space:
see what are the effects of changing resource limits.

“Don't guess, measure,” but you also have to measure something meaningful...

Our first major optimization, and perhaps the easiest to do, is the batch size.



The code is a little too large for the slides so let's go over it in a code editor.

The bert-large-uncased model is about 340 MB.

It's uncased because it makes no distinction between capitals and lower-case letters, e.g., it sees “Word” and “word” as equivalent.

```
jzarnett@ecetesla0:~/github/ece459/lectures/live-coding/L24$ python3 dummy_data.py
Starting up. Initial GPU utilization:
GPU memory occupied: 0 MB.
Initialized Torch; current GPU utilization:
GPU memory occupied: 417 MB.
Some weights of BertForSequenceClassification were not initialized
from the model checkpoint at bert-large-uncased and are newly initialized:
['classifier.bias', 'classifier.weight']
You should probably TRAIN this model on a down-stream task to be able
to use it for predictions and inference.
GPU memory occupied: 1705 MB.
torch.cuda.OutOfMemoryError: CUDA out of memory. Tried to allocate 20.00 MiB (GPU 0;
7.43 GiB total capacity; 6.90 GiB already allocated; 16.81 MiB free; 6.90 GiB
reserved in total by PyTorch) If reserved memory is >> allocated memory try setting
max_split_size_mb to avoid fragmentation. See documentation for Memory Management
and PYTORCH_CUDA_ALLOC_CONF
```

Let's See the Problem

I asked nvidia-smi...

+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									
NVIDIA-SMI		470.199.02		Driver Version: 470.199.02			CUDA Version: 11.4		
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									
GPU	Name	Persistence-M		Bus-Id	Disp.A	Volatile	Uncorr.	ECC	
Fan	Temp	Perf	Pwr:Usage/Cap	Memory-Usage		GPU-Util	Compute M.		
						MIG M.			
=====+=====+=====+=====+=====+=====+=====+=====+=====+=====									
0	Tesla P4	Off		00000000:17:00.0	Off			0	
N/A	42C	P0	23W / 75W	0MiB / 7611MiB		1%	Default		
						N/A			
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									
Processes:									
GPU	GI	CI	PID	Type	Process name	GPU Memory			
		ID	ID	Usage					
=====+=====+=====+=====+=====+=====+=====+=====+=====+=====									
No running processes found									
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									

7 611 MB is not enough for this model...



Problem is we don't have any bigger VRAM cards.

What I actually did next was change to a smaller version of the model, bert-base-uncased.

It's significantly smaller (110 MB) and something the card could handle.

```
jzarnett@ecetesla0:~/github/ece459/lectures/live-coding/L24$ python3 dummy_data.py
Starting up. Initial GPU utilization:
GPU memory occupied: 0 MB.
Initialized Torch; current GPU utilization:
GPU memory occupied: 417 MB.
Some weights of BertForSequenceClassification were not initialized from the
model checkpoint at bert-base-uncased and are newly initialized:
['classifier.weight', 'classifier.bias']
You should probably TRAIN this model on a down-stream task to be able to use
it for predictions and inference.
GPU memory occupied: 887 MB.
{'loss': 0.0028, 'learning_rate': 1.1718750000000001e-06, 'epoch': 0.98}
{'train_runtime': 109.6152, 'train_samples_per_second': 4.671,
'train_steps_per_second': 4.671, 'train_loss': 0.0027378778694355788, 'epoch': 1.0}
Time: 109.62
Samples/second: 4.67
GPU memory occupied: 3281 MB.
```

Then I needed to experiment some more with batch size to find the ideal.

Batch Size	Time (s)	Samples/s	Memory Used (MB)	Utilization (%)
1	109.62	4.67	3 281	43.1
2	85.82	5.97	3 391	44.6
4	72.18	7.09	4 613	60.6
8	66.70	7.68	7 069	92.9

Unsurprisingly, choosing batch size of 9 leads to OOM.

We seem to be memory limited – let's see what we can do?

First idea: **gradient accumulation**.

Calculate gradients in small increments rather than for the whole batch.

Experimenting with Gradient Accumulation

Steps	Time (s)	Samples/s	Memory Used (MB)	Utilization (%)
1	66.06	7.75	7 069	92.9
2	63.96	8.01	7 509	98.7
4	62.81	8.15	7 509	98.7
8	62.65	8.17	7 509	98.7
16	62.42	8.20	7 509	98.7
32	62.44	8.20	7 509	98.7
128	62.20	8.23	6 637	87.2
1024	61.78	8.29	6 637	87.2
4096	62.16	8.24	6 637	87.2

I got suspicious about the 128 dropoff in memory usage and it made me think about other indicators – is it getting worse somehow?


The output talks about training loss...

Gradient Accumulation Steps	Loss
1	0.029
2	0.070
4	0.163
8	0.169
16	0.447
32	0.445
128	0.435
1024	0.463
4096	0.014

Is that concerning?

We need to do some validation of how it's going...

We'll train and validate using some Yelp data.



The python code here is significantly different but we'll take a look.

I've skipped some intermediate results since at 9 minutes to calculate it takes a while to fill in all the values above.

GA Steps	Time (s)	Samples/s	Memory Used (MB)	Final Accuracy
1	538.37	5.56	7 069	0.621
8	501.89	5.98	7 509	0.554
32	429.70	6.98	7 509	0.347
1024	513.17	5.85	7 509	0.222

Interesting results – maybe a little concerning?

Increasing the gradient accumulation does change the effective batch size...

Batch size too large may mean less ability to generalize (overfitting).

But smaller isn't always better either; can underfit the data.

In the Yelp example, I get worse accuracy with batch size of 1 than 4, and 4 is worse than 8. There really is no magic number.

Gradient Checkpointing: increase compute time to save memory.

Instead of saving activations, save only some, recompute the rest.

Does it work?

Trying this out with batch size of 8, the total time goes from 66.70 to 93.07s and the memory from 7 069 down to 3 619 MB.

As expected, we got slower but used less memory.

Maybe it means we can increase the batch size?

Raising it to 16 means the time was 100.55s but still only 3 731 MB

Increasing the batch size a lot to finish faster might work eventually.

And no, using this checkpointing even with a batch size of 1 is not sufficient to run the bert - large - uncased model on ec2es1a0.

And remember that excessively large batch sizes aren't good...

Mixed Precision: Use less accurate types (16 vs 32-bit).

Data Preloading: Multithread to get data to GPU faster.

Other ideas: Mixture of Experts, bigger GPU, multiple-GPU...



- Accuracy for time?
- Memory for CPU?
- Err on the side of over- or under-fitting?

In the next few years, technologies and decision-making for this will become much more sophisticated.

But in the meantime, we can have a lot of fun experimenting and learning.

